

ML-BASED COMBINATORIAL OPTIMIZATION FOR SOLVING THE REALISTIC VEHICLE ROUTING PROBLEM

Final Presentation
Student Research Project, Winter Semester 22/23
02.12.2022

Project Supervisor: Daniela Thyssens

Student team: Ashish Sahu, Ilayda Topuz,
Jaishree Janu, Jatin Garg, Sophia Lawal




AGENDA

- — Introduction to Vehicle Routing Problem (VRP)
- — Research Workflow
- — Methodology
 - Realistic data generation
 - Setup baseline model and training environment
 - Proposed model architecture for time dependent dynamism
 - Experiments with various model configurations
- — Experiments
 - Experimental set-up
 - Evaluation of test results vs baseline
 - Training curves
 - Sample tour solutions
- — Conclusion and Discussion



INTRODUCTION TO VRP (OVERVIEW)



Graph Structure

Nodes: Co-ordinates, Demands
Edges: Distance between node-pairs

- All vehicles start from the depot node
- Fulfill all N customer demands
- Visit all customers once (no partial fulfilment)
- Final stop for all \rightarrow Depot
- Time window constraints have not been considered



PROBLEM STATEMENT

RESEARCH WORKFLOW

Realistic Data Generation

- Extracted co-ordinates from Open Street Map (OSM)
- Generated graph instances by sampling nodes and their respective distances
- Extracted live traffic variations to generate time dependent dynamism in graph edge features

Baseline implementation

- Setup training components and model architecture in Tianshou & PyTorch frameworks.
- Trained baseline model on the generated realistic datasets (also with edge dynamism)
- Evaluated the results for test data and extract optimal tours, then compare with ours

Graph Attention with edge feature encoding

- Experimented with different model configurations to encode edge features in graph context
- Finalized the model architecture for feeding edges to the attention-based encoder
- Experimented and evaluated the model to compare with baselines

Experiments

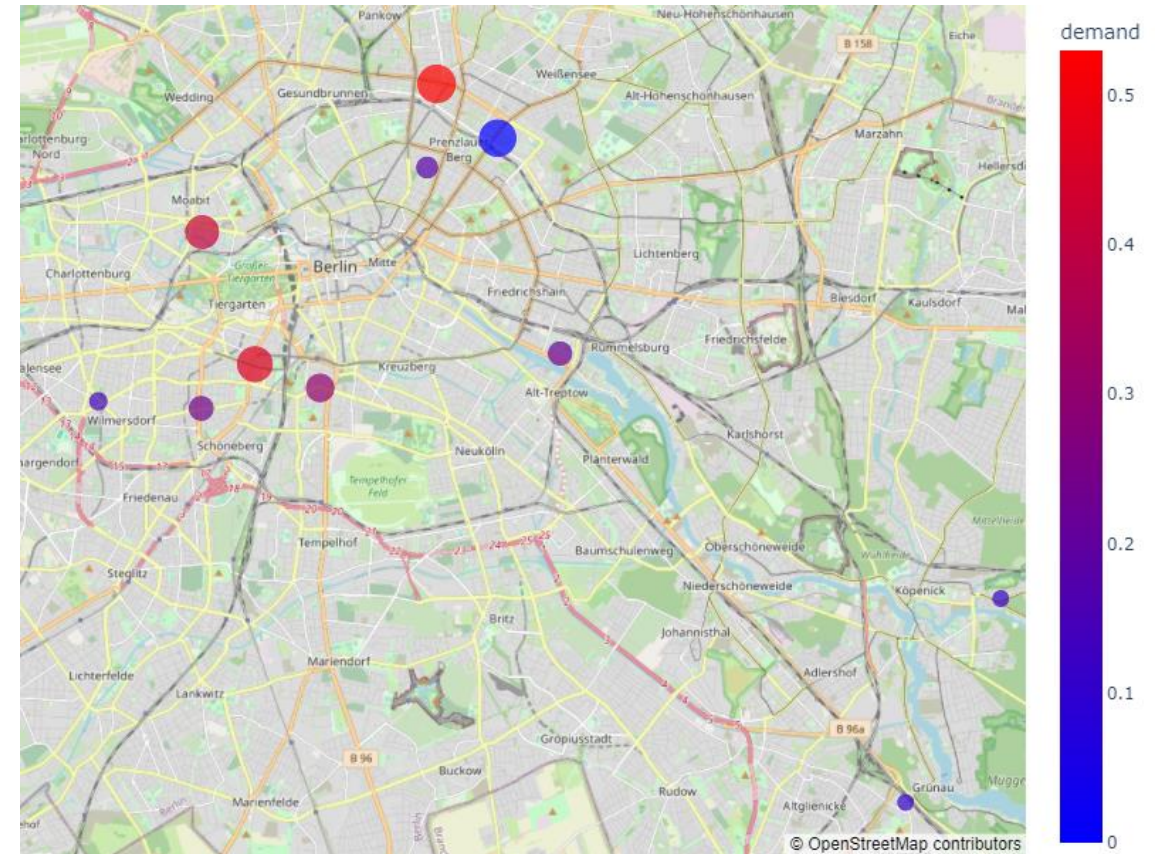
- Evaluation Metrics → Rewards and Loss vs trained epochs
- Performed hyperparameter tuning to optimize model training process
- Performed various experiments with different graph sizes and datasets with dynamic as well as static distance edges.

REALISTIC DATA GENERATION

GENERATE GRAPH INSTANCES

Real-time locations from Open Street Map (stores in Berlin)

- For a given address query [city_name, shop_type], coordinates and distances are taken from OSM*
- OSM gives real and asymmetric distances between respective couple of coordinates. Demands are generated stochastically from random Gaussian distributions.
- Efficiently generated large scale realistic graphs → instead of naively sampling with iterator, we created a large pool of co-ordinates and global distance matrix then subsampled graph instances.
- 2000 locations of supermarkets, convenience stores & grocery shops in Berlin was taken for generating global node pool.




*Accessed Overpass API for OSM node coordinates and distance features. <https://python-overpy.readthedocs.io/en/latest/introduction.html>

REALISTIC DATA GENERATION

DYNAMIC EDGES BY TRAFFIC INDEXING

- For more realistic graph instances, we studied the traffic congestion levels in the city (Berlin) and created timing distortion masks
- These masks are in the form of adjacency matrices, containing delay level information for augmenting graphs with delay-adjusted edge features
- Downloaded hourly traffic variation of past three years and performed aggregations (AVG.)
- A 30% congestion level means that on average, travel times were 30% longer than during the static conditions.
[This means that a 30-minute trip driven in free-flow condition will take 9 minutes longer when the congestion level is at 30%.]




MODEL IMPLEMENTATION

TRAINING ENVIRONMENT SETUP

What is current state?

x_t^i : current position of the vehicle
 o_t : remaining capacity of the vehicle
 M_t : action mask: possible actions
at current time step

Get actions from the model



Make environment
step as per the
selected actions

r_t^{ij} : action rewards

Set-up graph environment

n^i : node features from the graph
 e_{ij} : edge: distance between nodes
 d^i : customer demands at the node

MODEL IMPLEMENTATION

BASELINE MODEL ARCHITECTURE

Encoder-Decoder framework for solving VRP

ENCODER

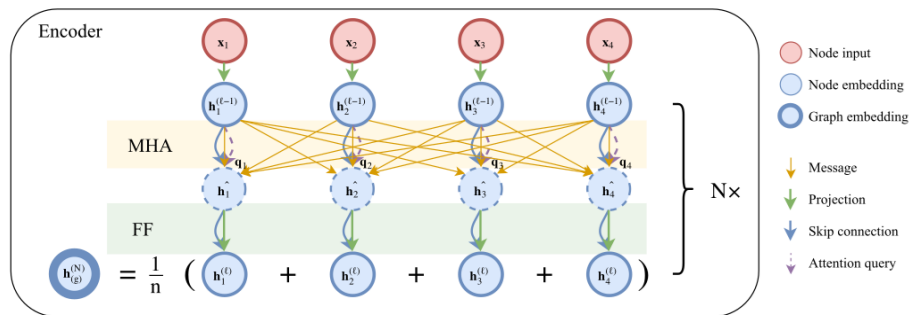


Fig. 2 : Multi-head Attention based encoder inspired from Transformer Architecture

Takes raw node embeddings and generate fixed size graph context

DECODER

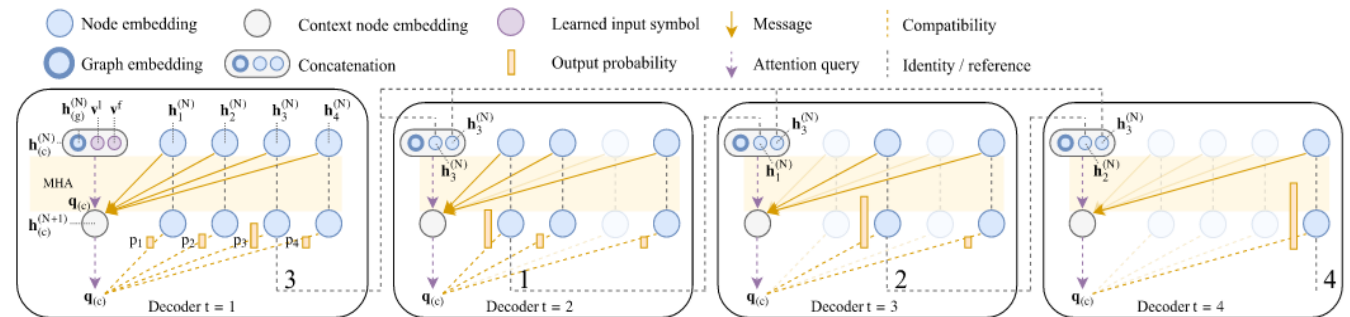
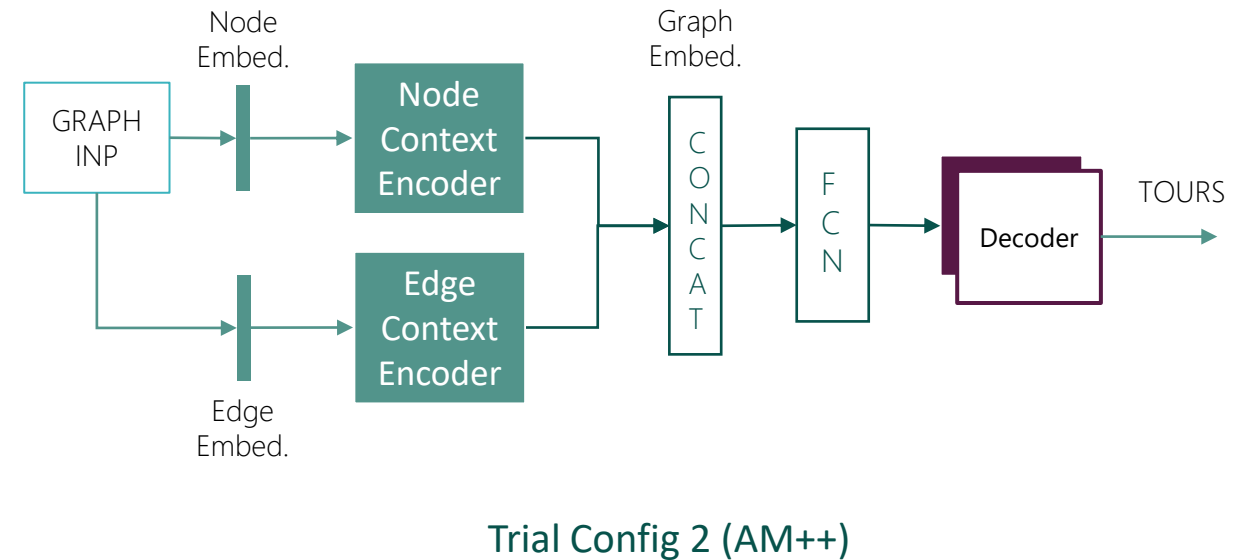
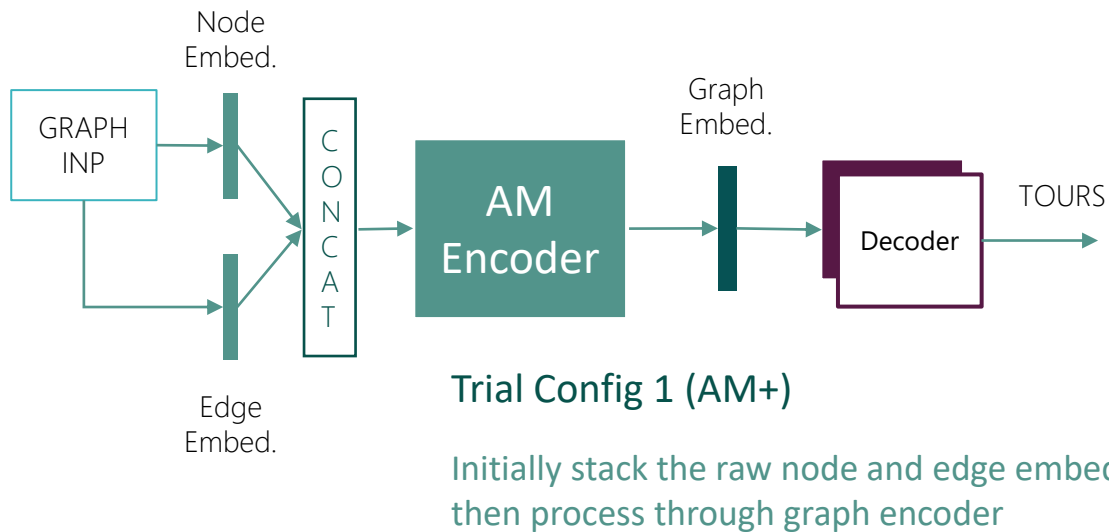
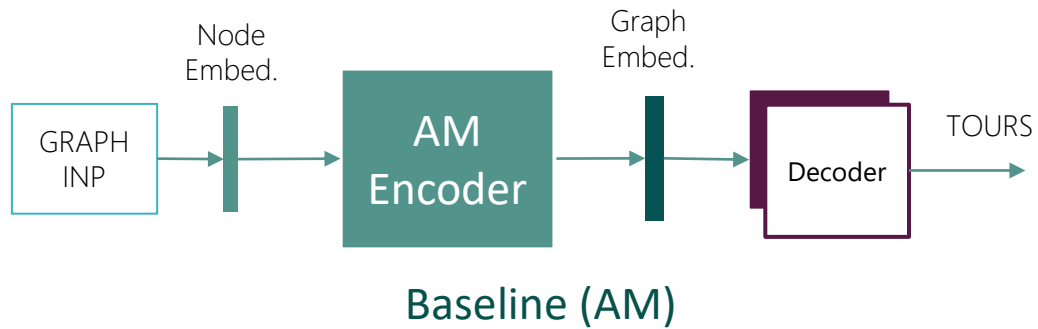


Fig. 3 : Attention based decoder

Takes graph context embeddings and generate tours autoregressively

MODEL IMPLEMENTATION

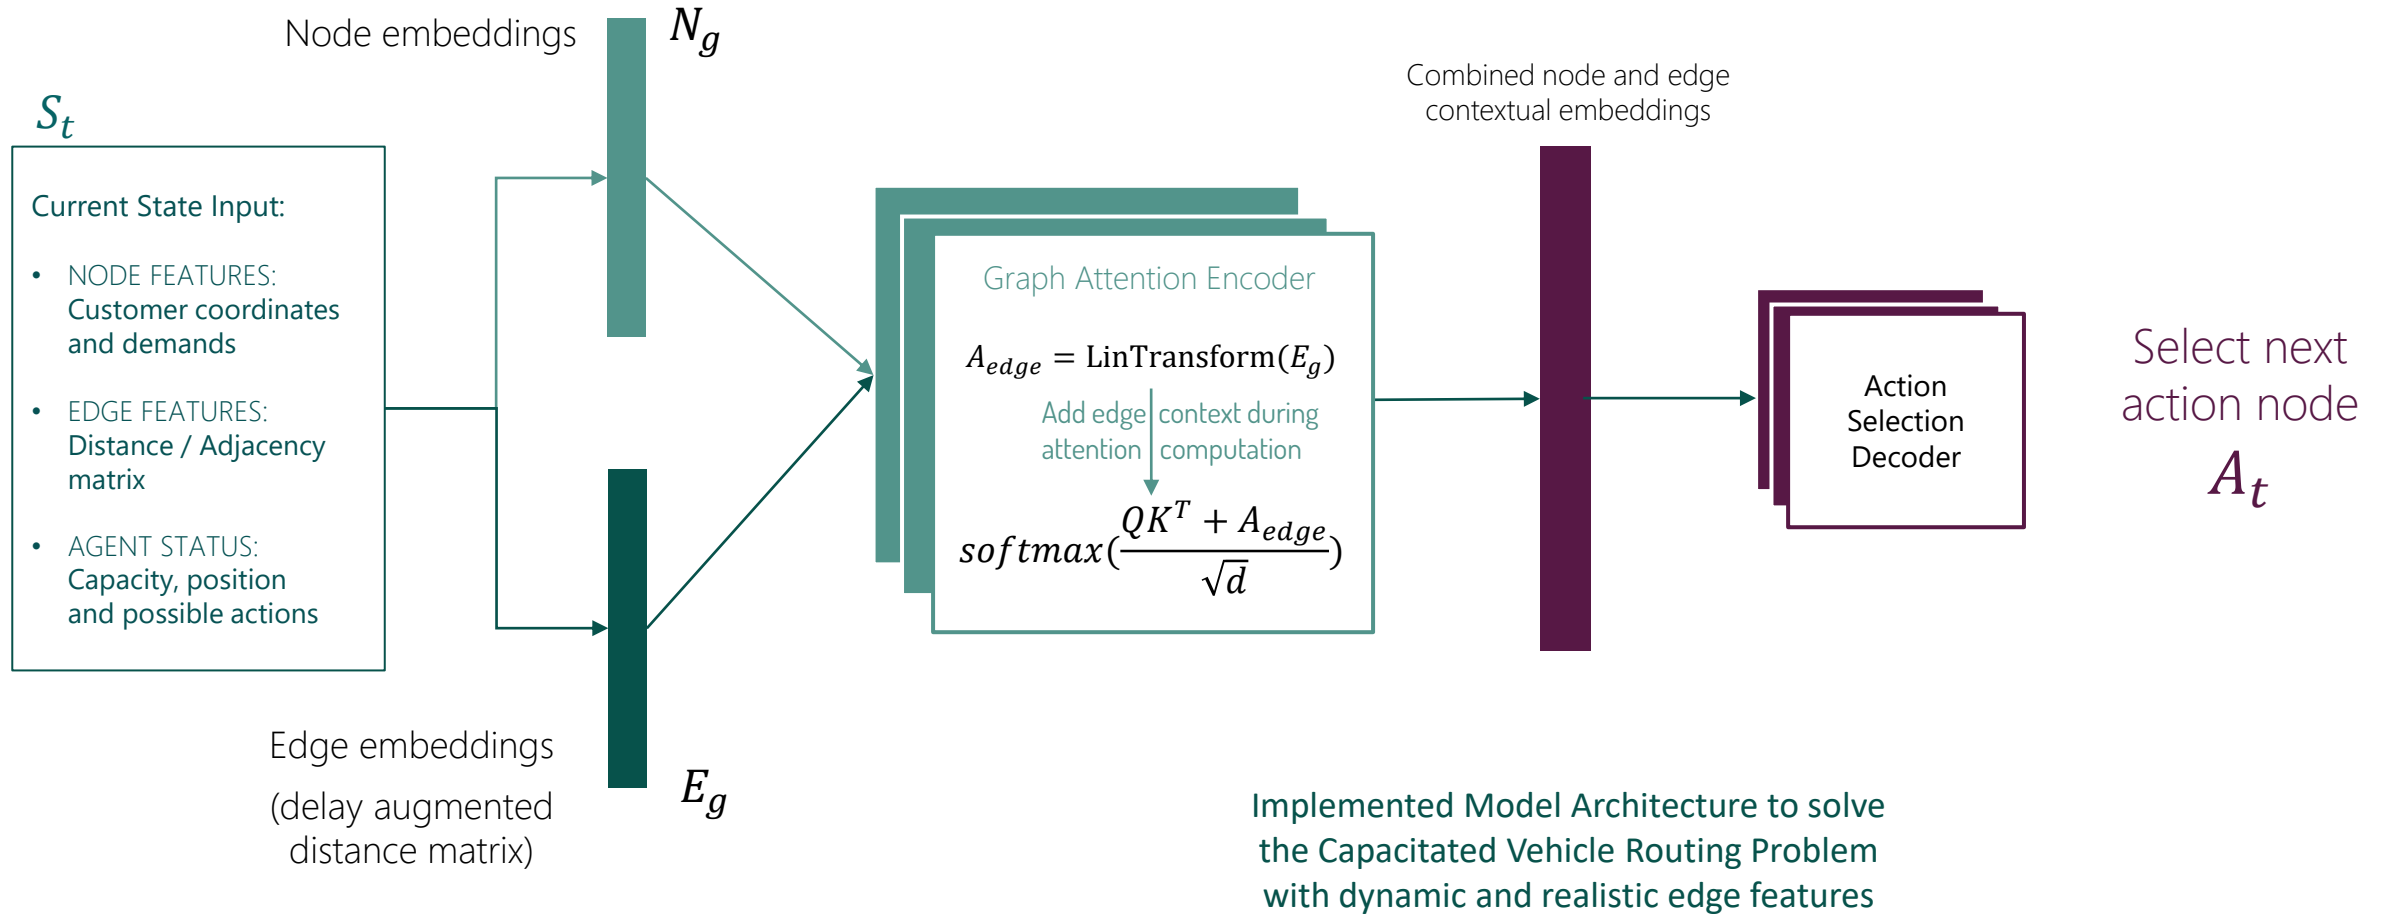
EXTENDING BASELINE WITH EDGE CONTEXT



Generate node and edge contexts, then stack and bring to same embedding space with FCN

MODEL IMPLEMENTATION

ENCODING EDGES FOR GRAPH ATTENTION



EXPERIMENT PROCESS

DATASETS, GRAPH SIZES AND PARAMETER TUNING

- Took 2000 locations of stores in Berlin (from Open Street Map), determined all distance pairs to create a global distance matrix
- One dataset with time-indexed edged features (*RealData-D*) while other with static node distances (*RealData-S*)
- Each graph instance was augmented with 24 delay levels, representing congestion levels throughout the day
- Different graph sizes: 10, 20, 30 and 50
- Training size: 10,000; Test size: 5,000. Also generated large scale datasets with 50k, 100k & 500k instances.
- Grid search for tuning hyperparameters (learning rate, weight decay, etc.)

EXPERIMENT PROCESS

EVALUATION OF RESULTS

Dataset	Realistic Data (static edges)			Realistic Data (time-indexed edges)		
Graph size	Attention Model (AM)	Edge+ AM (Ours)	Gap (vs base)	Attention Model (AM)	Edge+ AM (Ours)	
10	81.576	<u>81.114</u>	-0.56%	106.312	<u>105.924</u>	-0.36%
20	170.821	<u>150.645</u>	-11.81%	201.274	<u>200.606</u>	-0.33%
30	<u>252.196</u>	254.413	+0.88%	<u>309.903</u>	311.544	+0.53%

Best Test Reward values (tour length in km)


Total data size: 20k static, 24k dynamic graphs

Test / Train ratio: 50%


Epochs trained: 30

EXPERIMENT TRAINING CURVES AND REWARDS

Training loss vs Epochs



Collected train rewards vs Epochs

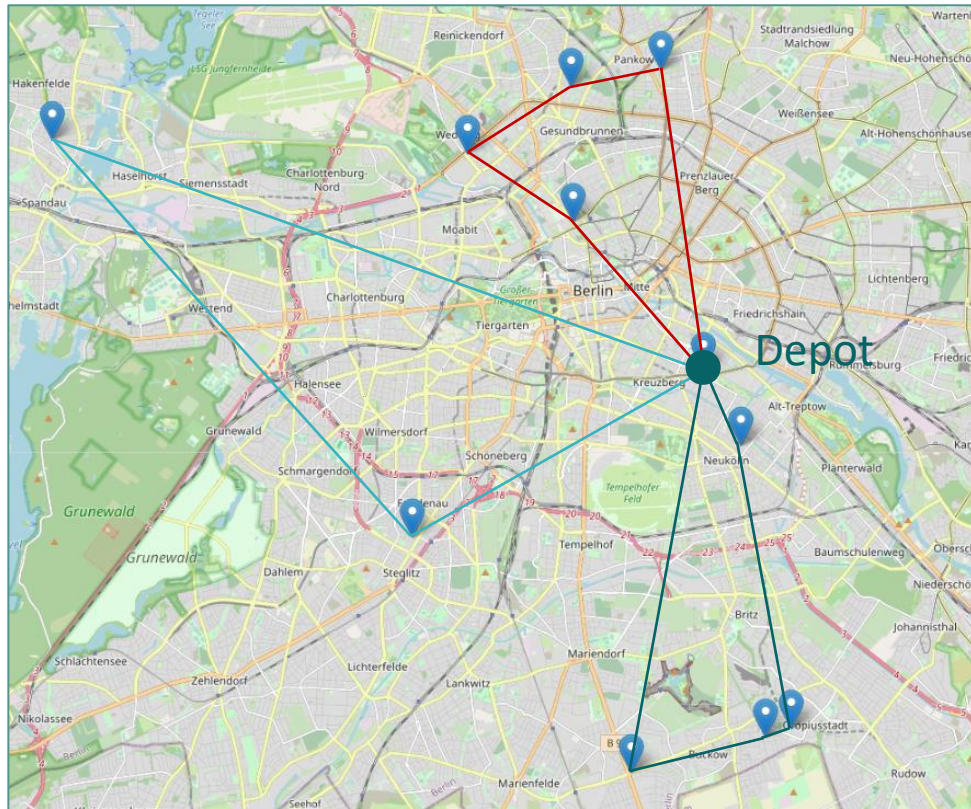


Training setting: Realistic static edge feature data
Total data size: 20k static, 24k dynamic graphs
Test / Train ratio: 50%
Epochs trained: 30

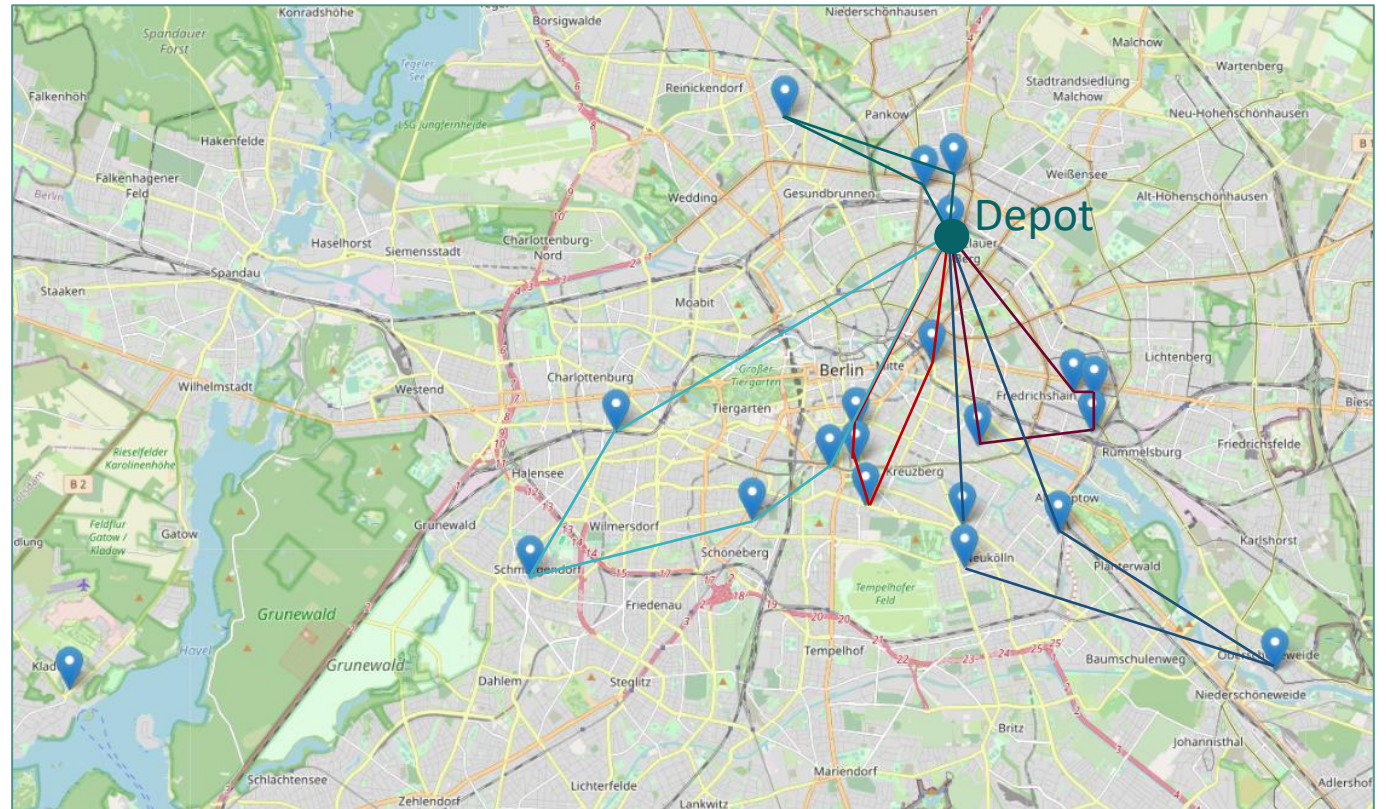
EXPERIMENT PROCESS

SAMPLE SOLUTION REALISTIC TOURS

N = 10; Length = 82.754 Km



N = 20; Length = 202.916 Km




CONCLUSION AND FUTURE SCOPE

- In this project, we have solved the Capacitated Vehicle Routing Problem (CVRP) by Deep Reinforcement Learning on realistically generated datasets.
- For both the datasets, our approach achieves better performance for small and medium graph sizes (in the given training setting). For $N=30$, the method is performing at par with the Attention Model baseline
- We have generated large scale realistic datasets for CVRP problems. Our method is scalable and can be applied to different cities to adapt routing solutions as per the traffic congestion in the city/state.
- As a future step, we will scale up the model training with large scale data (500k, 100k, etc.) and will try to cope with OOM limitations with computational tricks. Furthermore, we can evaluate our realistic generation method on various ML-based solutions – can it act as a benchmark generation method for the community?

REFERENCES

- [1] Irwan Bello, Hieu Pham, Quoc V. Le, Mohammad Norouzi, and Samy Bengio. Neural Combinatorial Optimization with Reinforcement Learning. CoRR, abs/1611.09940, 2016. arXiv: 1611.09940.
- [2] Yoshua Bengio, Andrea Lodi, and Antoine Prouvost. Machine learning for combinatorial optimization: A methodological tour d’horizon. European Journal of Operational Research, 290(2):405–421, April 2021.
- [3] Lu Duan, Yang Zhan, Haoyuan Hu, Yu Gong, Jiangwen Wei, Xiaodong Zhang, and Yinghui Xu. Efficiently Solving the Practical Vehicle Routing Problem: A Novel Joint Learning Approach. In Proceedings of the 26th ACM SIGKDD International Conference on Knowledge Discovery & Data Mining, KDD ’20, pages 3054–3063, New York, NY, USA, August 2020. Association for Computing Machinery.
- [4] Wouter Kool, Herke van Hoof, and Max Welling. Attention, Learn to Solve Routing Problems! Feb 2019.
- [5] Yeong-Dae Kwon, Jinho Choo, Byoungjip Kim, Iljoo Yoon, Seungjai Min, and Youngjune Gwon. POMO: Policy Optimization with Multiple Optima for Reinforcement Learning. Advances in Neural Information Processing Systems, pages 21188–21198, 2020. arXiv: 2010.16011.
- [6] Jingwen Li, Yining Ma, Ruize Gao, Zhiguang Cao, Andrew Lim, Wen Song, and Jie Zhang. Deep reinforcement learning for solving the heterogeneous capacitated vehicle routing problem. IEEE Transactions on Cybernetics, pages 1–14, 2021.
- [7] Yining Ma, Jingwen Li, Zhiguang Cao, Wen Song, Le Zhang, Zhenghua Chen, and Jing Tang. Learning to Iteratively Solve Routing Problems with Dual-Aspect Collaborative Transformer. May 2021.
- [8] Daniela Thyssens, Jonas Falkner, and Lars Schmidt-Thieme. Supervised permutation invariant networks for solving the cvrp with bounded fleet size. 2022.
- [9] Oriol Vinyals, Meire Fortunato, and Navdeep Jaitly. Pointer Networks. In Advances in Neural Information Processing Systems, volume 28. Curran Associates, Inc., 2015.
- [10] Veličković, P., Cucurull, G., Casanova, A., Romero, A., Liò, P. & Bengio, Y. (2017). Graph Attention Networks. 6th International Conference on Learning Representations.

THANK YOU FOR YOUR ATTENTION!



Let's
Discuss..



PROBLEM STATEMENT AND RESEARCH SCOPE

APPENDIX AI



GRAPH STRUCTURE

- One depot and multiple customer locations (N) with demands + co-ordinates
- Heterogeneously capacitated: Fleet of vehicles with different capacities (cars, trucks, etc.)
- Directed asymmetric graph: Distances between node-pairs are asymmetric, $d_{AB} \neq d_{BA}$
- Dynamic distances: Distances can change when vehicle in transit due to traffic/weather (TBD*)



LEARNING METHOD

- Dual objective: Minimize total distance travelled by each vehicle + number of vehicles used
- Construction approach: Generate policy regressively, at each step predicting next node distribution
- DeepRL methodology: Encoder-Decoder framework, determine near optimal solution



ASSUMPTIONS

- All vehicles start & end at depot node, and each customer is visited exactly once
- Vehicle has sufficient capacity to fulfill all customer demands within one tour-path
- Demands are fixed, all the vehicles have same speed and no time window constraints

RELATED WORK AND THEIR LIMITATIONS

APPENDIX A2

—• Pointer Networks •—

- RNN + Attention based methods to generate optimal tour policy with sequence-to-sequence learning approach
- Supervised Learning setting requires carefully generated ground-truth labels for good model predictions, only for TSP
- **Actor-Critic Algorithm:** PN without supervised setting, uses training instance tour costs as Monte-Carlo estimates for stochastic policy

[Bello et al, 2016; Vinyals et al. 2015]

—• Graph Neural Networks •—

- Graph Convolutions (GCN) and Graph Attention (GAT) based encoders can be used to learn node context embeddings at different time steps. Attention based decoders used to compute distribution over policies
[Duan et al, 2020; Kool et al, 2019]
- For HCVRP (heterogeneous capacity), a 2-level decoder was proposed: first one selects a vehicle, and the other predicts next node to visit for that vehicle
[Li et al, 2019]
- These methods use either Euclidean distance between stochastically generated nodes or homogeneous fleet for learning, failing to model complex real-world scenario

Training
Data
(Graph)

Encoder

Generates graph embeddings for all input nodes

Decoder

Predicts near-optimal tour policy from graph context

Generate low cost path for all vehicles

CONSTRUCTION



IMPROVEMENT

PROBLEM APPROACH

TRAINING OBJECTIVE AND ALGORITHM

LOSS FUNCTION

- Minimize total length of tours: sum of distances between selected nodes
- Learning objective function:
 - $m_{\pi_t \pi_{t+1}}$ is travel distance between current time position and selected next node

$$\min \sum_{t=1}^{T-1} m_{\pi_t \pi_{t+1}}$$

TRAINING ALGORITHM

- REINFORCE – Policy Gradient Algorithm (with baseline)
- Gradient estimator for parameter update:

$$\nabla \mathcal{L}(\theta|s) = \mathbb{E}_{p_{\theta}(\pi|s)} [(L(\pi) - b(s)) \nabla \log p_{\theta}(\pi|s)]$$

PROBLEM APPROACH

HOW TO PERFORM A STEP IN ENVIRONMENT

Transition

Given the current state, transit to next node with highest probability

Transition from previous state S_t to the next state S_{t+1} based on the performed action A_t :

Vehicle capacity update: $o_{t+1} = o_t - d_t^i$

Action mask update: $M_{t+1} = \text{mask}(M_t, A_t)$

Partial route update: $\pi_{t+1} = [\pi_t, A_t]$

Current position update: $x_{t+1}^i = A_t$

The agent will be trained to maximize reward

Reward

Reward is defined as the negative of the tour length, which corresponds to the distance between nodes (from edge feature matrix).

RESEARCH IDEA

DYNAMIC DISTANCES BETWEEN NODES



Distance travelled by a vehicle can change due to weather/construction, and so the tour length

Can we enable the model to handle this dynamism in edge features during the training time?




$$d^w = d (\pm \Delta d)$$

There can be delays due to weather, traffic condition, etc. any time



$$d^a = d (\pm \Delta d)$$

There might be some alternate routes with slightly deviated distance/time



How to create this dynamism in graph, model and data generation?

REALISTIC DATA GENERATION

DYNAMIC DISTANCING BETWEEN NODES

There can be delays due to weather, traffic condition, etc. any time



$$d^w = d (\pm \Delta d)$$




On some days, shortest road might be fully / partially closed



$$d^a = d (\pm \Delta d)$$


There might be some alternate routes with slightly deviated distance/time



Edge features

Static distances between the nodes


Obtained from Open Street Map



Timing distortions

Possible +/- changes in distance/time

Noise distribution by domain knowledge



Noise Weights

Learned model weights for dynamism

Fine-tune timing distortions by weights